

Topic 13

Heap, Set and Map

資料結構與程式設計
Data Structure and Programming

11/25/2015

Linear Data Types

- ◆ In previous topic and Homework #5, we have learned linear data types like list and array
 - Tradeoffs between insert/delete/find operators
 - Memory overhead
 - ➔ Constant time for “push_back()” or “push_front()” operation
- ◆ The best way to use linear data types is ---
 - Data are recorded in a linear sequence (i.e. only push_back or push_front is needed)
 - Linearly traverse each element (i.e. for(...; li++))
 - No “find”, “insert any”, nor “delete any”

Consider the Scenario...

- ◆ Suppose we are assigning jobs sequentially to several machines ---
 - One job to one machine and we record the accumulated runtime for each machine.
 - Our machine selection criteria is to “even out” the runtime of the machines.
 - In other words, we would like to pick the machine with least accumulated runtime for the next job
 - ➔ Do we need to sort ALL the elements?
 - ➔ Need a priority queue

Priority Queue


- ◆ An ADT that supports 2 operations
 - Insert
 - Delete min(or max)
- ◆ An element with arbitrary priority can be inserted to the queue
- ◆ At any time, it should take constant time to find the element with min(or max) priority and remove it from the list
 - Need to figure out which is the one with next lowest(highest) priority efficiently

Using List or Array?



- ◆ Use linear ADT with an extra field to record the element with min(max) priority
 - Insert: $O(1)$
 - Delete min(max): $O(n)$ (why?)
- ◆ As we learn before, $O(n)$ is not good. We would prefer an ADT with $O(\log n)$ for both operations

Min (Max) Heap

- ◆ A complete binary tree in which the key value in each node is no larger (smaller) than its children



Remember that we can use array to implement a complete binary tree...




◆ Parent
= child / 2

◆ Child
= Parent * 2
or Parent * 2 + 1

MinHeap Insertion

```
// Let n be the index of the last element
void MinHeap::insert(const T& x)
{
    int t = ++n; // next to the last
    while (t > 1) {
        int p = t / 2;
        if (x._key >= _heap[p]._key)
            break;
        _heap[t] = _heap[p];
        t = p;
    }
    _heap[t] = x;
}
```




What's the time complexity?

Delete Min Element

```

T& MinHeap::deleteMin()
{
    T ret = _heap[1];
    int p = 1, t = 2 * p;
    while (t <= n) {
        if (t < n) // has right child
            if (_heap[t]._key > _heap[t+1]._key)
                ++t; // to the smaller child
        if (_heap[n]._key < _heap[t]._key)
            break;
        _heap[p] = _heap[t];
        p = t;
        t = 2 * p;
    }
    _heap[p] = _heap[n--];
    return ret;
}

```




What's the time complexity?

Min(Max) Heap

- ◆ Simple implementation (just an array)
- ◆ Good insertion and deleteMin complexity
 - $O(\log n)$ vs. $O(n)$

What if you want to
delete min AND delete max?

Min-Max Heap




- Insert, delete min, delete max: all $O(\log n)$ (why?)

Deap

- ◆ Double-ended heap
 1. The root contains no element
 2. The left subtree is a min heap
 3. The right subtree is a max heap
 4. Let i be any node in the left subtree. Let j be the corresponding node in the right subtree. If such a j node does not exist, then let j be the corresponding parent of i .
 ➔ The key in node i is less than or equal to that in j .

Deap Example




- Insert, delete min, delete max: all $O(\log n)$ (why?)
 - But faster than min-max heap by a constant factor
 - Algorithm is simpler

Data Structure and Programming

Prof. Chung-Yang (Ric) Huang

13

Deap Implementation



- Given a node 'i', how to find the "corresponding parent" or "corresponding child"?
- When insertion or deletion, what should we do when the node value is greater/smaller than its corresponding parent/child?

Data Structure and Programming


Prof. Chung-Yang (Ric) Huang

14

More Varieties of Heaps: Leftist Heap

- ◆ In contrast to a *binary heap*, a leftist heap attempts to be very unbalanced.

- s-value(v): the distance to the nearest leaf.
- In addition to the heap property, the right child of each node has the lower s-value.



- ◆ Support "combine(heap1, heap2)" in $O(\log n)$

Data Structure and Programming


Prof. Chung-Yang (Ric) Huang

15

Leftist Heap: Huh?

- ◆ Remember: "combine(heap1, heap2)" in $O(\log n)$
 - Both "insert" and "deleteMin" operations can be realized by "combine". (How?)

```
combine(h1, h2) {
    compare(min(h1), min(h2));
    // let min(h1) < min(h2)
    if (right(h1) == NULL)
        right(h1) = h2;
    else
        combine(right(h1), h2);
    // h2 is now the combined heap
    if (s(right(h2)) > s(left(h2)))
        swap(right(h2), left(h2));
}
```




Data Structure and Programming

Prof. Chung-Yang (Ric) Huang

16

Leftist Tree: Combine

[src] <http://blog.yam.com/rockmanray/article/44962825>




Data Structure and Programming

Prof. Chung-Yang (Ric) Huang

17

More Varieties of Heaps: Binomial heap

- ◆ Binomial tree of order k
 - Binomial tree of order 0 is a single node
 - The root of a binomial tree of order k has k children, who are roots of binomial trees of order $k-1, k-2, \dots, 0$
 - Has exactly 2^k nodes; height = k
- ◆ Binomial heap
 - A collection of Binomial trees
 - Most operations have the complexity $O(\log n)$
 - But the amortized complexity is either $O(1)$ or $O(\log n)$



Data Structure and Programming

Prof. Chung-Yang (Ric) Huang

18

Binomial Heap: Properties

- ◆ Given a binomial heap with n nodes:
 - The node containing the min element is a root of B_0, B_1, \dots , or B_k .
 - It contains the binomial tree B_i iff $b_i = 1$, where $b_k \cdot b_{k-1} \cdot b_{k-2} \cdot \dots \cdot b_1 \cdot b_0$ is binary representation of n .
 - It has $\leq \lfloor \log_2 n \rfloor + 1$ binomial trees.
 - Its height $\leq \lfloor \log_2 n \rfloor$.

[src] <http://www.cs.princeton.edu/~wayn/leilei/ncg4andcs/pdf/BinomialHeaps.pdf>

Data Structure and Programming

Prof. Chung-Yang (Ric) Huang

19

Binomial Heap: Operations

- ◆ Similar to Leftist Heap, the operations of Binomial Heap can be realized by the “compose” (aka. “meld”) operation.
- ◆ Compose operation:
 - Binary addition
 - Given two binomial heaps

$$H_1 := \{ (B_3, B_2, B_1, B_0) = (1, 1, 0, 1) \}, \text{ and}$$

$$H_2 := \{ (B_4, B_3, B_2, B_1, B_0) = (1, 0, 1, 0, 1) \}.$$
 The composed binomial heap

$$H_m := \{ (B_5, B_4, B_3, B_2, B_1, B_0) = (1, 0, 0, 0, 1, 0) \}.$$

Data Structure and Programming

Prof. Chung-Yang (Ric) Huang

20

Binomial Heap: Compose Operation


- ◆ Atomic operation:
 - Given two binomial trees B_i, B_j , with the same order k , then $\text{compose}(B_i, B_j)$:
 1. Connect the roots r_i, r_j of B_i, B_j .
 2. Choose $\min(r_i, r_j)$ as the root of the composed tree
 3. The composed tree is of order $k+1$
 → What if we have three binomial trees with the same order?
- ◆ The compose operation of two binomial heaps:
 1. Align the binomial trees of both heaps
 2. From the trees with the least order, perform tree composition
 3. Propagate to the next order of tree if necessary
- ◆ What's the time complexity? $O(\log n)$

Binomial Heap: Other Operations

- ◆ FindMin
 - // remember: It has $\leq \lfloor \log_2 n \rfloor + 1$ binomial trees
 - $O(\log n)$
- ◆ DeleteMin
 - Note: after the "min" is removed, the corresponding binomial tree (of order k) is broken and becomes k binomial trees
 - It just becomes "compose" operations of some binomial trees // How many?
 - $O(\log n)$
- ◆ DeleteNode(iterator pos)
 - $O(\log n)$
- ◆ Insert(x)
 - $O(\log n)$

More Varieties of Heaps: Fibonacci heap

- ◆ Fibonacci heap
 - Especially useful when $\text{deleteMin}()$ & $\text{delete}(n)$ are rarely called → amortized $O(\log n)$
 - All other operations are $O(1)$




Fibonacci Heap

- ◆ Basic idea
 - Similar to binomial heaps, but less rigid structure
 - Binomial heap: eagerly consolidate trees after each insert (maintain binomial structure)
 - Fibonacci heap: lazily defer consolidation until next **delete-min**
- ◆ Properties
 - Set of heap-ordered trees.
 - Maintain pointer to minimum element
 - Set of marked nodes

Fibonacci Heap: Insert Operation

- ◆ Create a new singleton tree.
- ◆ Add to root list; update min pointer (if necessary) $\rightarrow O(1)$



(Ref) <https://www.cs.princeton.edu/~wajwe/teaching/fibonacci-heap.pdf>

Fibonacci Heap: DeleteMin Operation

- ◆ Let H be a Fibonacci heap and x be a node
 - Rank(x): number of children of node x
 - Rank(H): max rank of any node in heap H
 - Tree(H): number of trees in heap H
 - ◆ DeleteMin
 - Delete min; meld its children into root list; update min
 - Consolidate trees so that no two roots have same rank
- \rightarrow Time complexity: $O(\text{rank}(H)) + O(\text{trees}(H))$
 \rightarrow Amortized cost: $O(\text{rank}(H))$

Heap Operations Supported in STL

- ◆ STL does not have a "heap" class
 - Instead, it support several operations that can operate on "array" like data structure
 - ◆ Operations
 - void make_heap(first, last[, comp]);
 - void push_heap(first, last[, comp]);
 - void pop_heap(first, last[, comp]);
 - void sort_heap(first, last[, comp]);
 - bool is_heap(first, last[, comp]);
- \rightarrow first, last: RandomAccessIterator
 \rightarrow comp: StrictWeakOrdering (optional)

Summary: Heap Structures

- ◆ Pros:
 1. Good complexity of "insert", "delete min(max)", ... operations
 2. Simple data structure (low memory overhead)
 3. Simpler algorithms (than BST)
- ◆ Con
 1. Data are not sorted
 - \rightarrow Still have $O(n)$ for "find" operation

Review: Binary Search Trees


- ◆ Binary Search Trees (BSTs)
 - Left subtree \leq this \leq right subtree
 - Complexity depends on the height of the tree
 - Worst case: can be degenerated as a tree with height $O(n)$
- ◆ Balanced BSTs
 - The heights of left subtree and right subtree are somewhat balanced
 - Height $\sim O(\log n)$
 - Examples: AVL, 2-3, 2-3-4, red-black, splay trees
 - Algorithms for their operations are complicated

Sorted ADT in STL

- ◆ Also classified as “Associative Containers”
 1. set
 2. multiset
 3. map
 4. multimap
- ➔ Implemented in “red black tree”

Red Black Tree

- ◆ A node is either red or **black**. The root is **black**
- ◆ All leaves are black (i.e. All leaves are same color as the root.)
- ◆ Every red node must have two **black** child nodes.
- ◆ Every path from a given node to any of its descendant leaves contains the same number of **black** nodes.
- ◆ Memory efficient
- ◆ Although balancing is NOT perfect, $O(\log n)$ for insert, delete, and find



class set in STL

- ◆ To store elements in a set
 - e.g. { 2, 3, 5, 7, 9 }
- ◆ `set<Key[, Compare, Alloc]>`
 - class Key: element type
 - class Compare: how the elements are compared (optional; default = `less<Key>`)
 - class Alloc: used for internal memory management (optional; default = `alloc`)

Member Functions in class set

1. `iterator begin() const;`
`iterator end() const;`
2. `pair<iterator, bool> insert(const value_type& x);`
`iterator insert(iterator pos, const value_type& x);`
`void insert(InputIterator, InputIterator);`
3. `void erase(iterator pos);`
`size_type erase(const key_type& k);`
`void erase(iterator first, iterator last);`
4. `iterator find(const key_type& k) const;`
5. `size_type count(const key_type& k) const;`
6. `iterator lower_bound(const key_type& k) const;`
`iterator upper_bound(const key_type& k) const;`
`pair<iterator, iterator> equal_range(const key_type& k) const;`

Other Functions for class set

1. `includes`
 - Check if one set is included in another
2. `set_union`
3. `set_intersection`
4. `set_difference`
5. `set_symmetric_difference`
 - $(A - B) \cup (B - A)$

class multiset in STL

- ◆ Unlike “set”, where elements with same value are stored only once, in multiset, they can be stored repeatedly
 - e.g. { 2, 3, 5, 5, 6, 7, 7, 7 }
- ◆ `multiset<Key[, Compare, Alloc]>`
 - class Key: element type
 - class Compare: how the elements are compared (optional; default = `less<Key>`)
 - class Alloc: used for internal memory management (optional; default = `alloc`)

class map in STL

- ◆ In many applications, data are associated with keys (or id's)
 - For example, (id, student record)
 - e.g. { (Mary, 90), (John, 85), (Sam, 71) ... }
- ◆ `class map<Key, Data[, Compare, Alloc]>`
 - class Key: compare data type
 - class Data: value type
 - class Compare: how the elements are compared (optional; default = `less<Key>`)
 - class Alloc: used for internal memory management (optional; default = `alloc`)

Example of using class map (1)

```
map<string, unsigned> scoreMap;
scoreMap["Mary"] = 90;
scoreMap["John"] = 85;
scoreMap["Sam"] = 71;
unsigned maryScore = scoreMap["Mary"];
cout << "Mary's score = " << maryScore << endl;
map<string, unsigned>::iterator mi;
mi = scoreMap.find("John");
if (mi != scoreMap.end())
    cout << "John's score = " << (*mi).second << endl;
→ How about "map<const char*, unsigned>"?
```

Comments about map::operator []

- ◆ Since operator[] might insert a new element into the map, it can't possibly be a const member function.
- ◆ Note that the definition of operator[] is extremely simple: m[k] is equivalent to `((m.insert(value_type(k, data_type()))).first).second`.
 - value_type = pair<Key, Data>
 - insert(value_type) returns a pair<map::iterator, bool>
- ◆ Strictly speaking, this member function is unnecessary: it exists only for convenience.

Bad example of using class map

```
map<const char*, unsigned> mmm;
map<const char*, unsigned>::iterator mi;
char buf[1024];
cin >> buf; mmm[buf] = 10;
cin >> buf; mmm[buf] = 20;
cin >> buf; unsigned s1 = mmm[buf];
cout << buf << " = " << s1 << endl;
cin >> buf; unsigned s2 = mmm[buf];
cout << buf << " = " << s2 << endl;
```

Example of using class map (2)

```
string str;
for (int i = 0; i < 5; ++i) {
    cin >> str; mm.insert(pair<string, int>(str, i));
}
while (1) {
    cin >> str;
    map<string, int>::iterator mi = mm.find(str);
    if (mi == mm.end()) {
        cout << "Not found!!" << endl;
        break;
    }
    cout << (*mi).first << " = " << (*mi).second << endl;
}
```

Conclusion: Set and Map

- ◆ “set” and “map” are useful data structures when we need to perform efficient “insert”, “erase”, and “find” operations
 - Usually implemented by balanced binary search trees
 - Implementation efforts can be high
 - Using STL may be a good choice
- ◆ Remember, unbalanced BSTs may not be a bad choice
 - Most randomly inserted BSTs are somewhat balanced
- ◆ Remember, there's no free lunch
 - Overhead in insert (vs. push_back)
 - If we don't need to do “erase” or “find” during insertions ... (what's the alternative?)